

Computer Architecture: Performance

Hossein Asadi (asadi@sharif.edu)

Department of Computer Engineering

Sharif University of Technology

Spring 2025



Copyright Notice

- Some Parts (text & figures) of this Lecture adopted from following:
 - Computer Organization & Design, The Hardware/Software Interface, 3rd Edition, by D. Patterson and J. Hennessey, MK publishing, 2005.
 - “Intro to Computer Architecture” handouts, by Prof. Hoe, CMU, Spring 2009.
 - “Computer Architecture & Engineering” handouts, by Prof. Kubiatowicz, UC Berkeley, Spring 2004.
 - “Intro to Computer Architecture” handouts, by Prof. Hoe, UWisc, Spring 2021.
 - “Computer Arch I” handouts, by Prof. Garzarán, UIUC, Spring 2009.
 - “Intro to Computer Organization” handouts, by Prof. Mahlke & Prof. Narayanasamy, Winter 2008.



Our Lectur Today



Topics Covered Today

- **Performance Metrics**
 - **Throughput**
 - **Latency**
 - **CPI**
- **Benchmarking**



Performance Metrics

- **Latency**
 - Time between start and finish of a single task
- **Throughput**
 - Number of tasks finished in a given unit of time
- Question:
 - $\text{Throughput} = 1 / \text{Latency} (?)$
 - $\text{Latency} = 1 / \text{Throughput} (?)$



Throughput vs. Latency

- Example: Airbus A380 vs. BAC Concorde
 - $\text{Speed}_{\text{Concorde}} > \text{Speed}_{\text{Boeing}}$
 - $\text{Latency}_{\text{Concorde}} < \text{Latency}_{\text{Boeing}}$
 - Throughput???
 - How to define throughput?





Airplane	Passenger Capacity	Speed (mph)
Airbus A380	850	610
Concorde	132	1350



Performance: Throughput vs. Latency

- Throughput = Passenger Capacity * Speed

Airplane	Pass. Capacity	Speed (mph)	Passenger Throughput
Boeing 747	470	620	291,400
Concorde	132	1350	178,200



Throughput vs. Latency: Single-Core Example

- Consider Two Single-Core Computers
 - **Less latency** (response time)
 - The one that finishes tasks faster
 - **More throughput**
 - The one that finishes more jobs
 - Or the one that finishes tasks faster
 - In this example
 - Latency ↓ ➔ Throughput ↑
 - Latency ↑ ➔ Throughput ↓
 - Not always true!



Throughput vs. Latency: Multi-Core Example

- Computer A
 - Two cores, running at 3Mhz
- Computer B
 - Four cores, running at 2.5Mhz
- $\text{Latency}_A < \text{Latency}_B$
- $\text{Throughput}_A < \text{Throughput}_B$

Computer	Frequency	# of cores
Type A	3 Mhz	2
Type B	2.5 Mhz	4



Intel 5th Generation Xeon Processors


2S PERFORMANCE GENERAL PURPOSE

SKU	CORES	BASE (GHz)	ALL CORE TURBO (GHz)	Max TURBO (GHz)	CACHE (MB)	TDP (Watts)	Maximum Scalability	DDR5 Memory Speed (1DPC)	UPI Links Enabled	Default DSA Devices	Default IAA Devices	Default QAT Devices	Default DLB Devices	Intel SGX Enclave Capacity (Per Processor)	Long life availability	Recommended Customer Pricing (RCP) in \$ US Dollars	Intel® On Demand Capable
8592+	64	1.9	2.9	3.9	320	350	2S	5600	4	1	1	1	1	512GB		\$11,600	✓
8580	60	2.0	2.9	4.0	300	350	2S	5600	4	1	0	0	0	512GB		\$10,710	✓
8570	56	2.1	3.0	4.0	300	350	2S	5600	4	1	0	0	0	512GB		\$9,595	✓
8568Y+	48	2.3	3.2	4.0	300	350	2S	5600	4	1	1	1	1	512GB		\$6,497	✓
8562Y+	32	2.8	3.8	4.1	60	300	2S	5600	3	1		1	1	512GB		\$5,945	✓
6548Y+	32	2.5	3.5	4.1	60	250	2S	5200	3	1	1	1	1	128GB	✓	\$3,726	✓
6542Y	24	2.9	3.6	4.1	60	250	2S	5200	3	1	0	0	0	128GB		\$2,878	✓
6544Y	16	3.6	4.1	4.1	45	270	2S	5200	3	1	0	0	0			\$3,622	✓
6526Y	16	2.8	3.5	3.9	37.5	195	2S	5200	3	1	0	0	0	128GB	✓	\$1,517	✓
6534	8	3.9	4.2	4.2	22.5	195	2S	4800	3	1	0	0	0	128GB		\$2,816	✓
5515+	8	3.2	3.6	4.1	22.5	165	2S	4800	3	1	1	1	1	128GB	✓	\$1,099	✓

© <https://wccfttech.com/intel-5th-gen-xeon-cpus-official-emerald-rapids-up-to-64-cores-320-mb-cache-prices/>



Throughput vs. Latency



Log-log plot of bandwidth and latency milestones



Performance Definition

- **Response Time**
 - Total time to complete a task (program)
 - Also called, **wall-clock time**, **elapsed time**
- Performance = 1 / response time
 - Response time $\downarrow \Rightarrow$ Performance \uparrow
 - $\text{Performance}(x) / \text{Performance}(y) =$
 $\text{Execution time}(y) / \text{Execution time}(x)$



Performance Definition (cont.)

- Response Time Consists of:
 - CPU time
 - CPU time spent on a program
 - I/O time
 - Time elapsed to wait for I/O transactions
- CPU Time
 - User CPU time
 - CPU time directly spent on a program
 - System CPU time
 - CPU time spent in OS doing tasks on behalf of a program



Performance Definition (cont.)

- System Performance
 - $1 / \text{Elapsed time}$
- CPU Performance
 - $1 / \text{User CPU time}$
- Clock Cycles
 - Clock periods, cycles, clock ticks, ticks, clocks



Response Time vs. Throughput

- Which One is More Important?
 - Response time or throughput
- **User** Perspective
 - **Response time** more visible
 - Unless a user runs bunch of tasks together
- **System Admin** Perspective
 - **Throughput** more visible



CPU Time

= CPU Clock Cycles * Clock Cycle Time

= CPU Clock Cycles / Clock Rate

- Example

- $\text{CPU}_A(\text{clock rate}) = 4 \text{ Ghz}$

- $\text{CPU}_B(\text{clock rate}) = 3 \text{ Ghz}$

- Which one runs program X faster?

- Depends on number of clock cycles spend on program X



CPU Time (cont.)

= CPU Clock Cycles * Clock Cycle Time

= CPU Clock Cycles / Clock Rate

- CPU Clock Cycles

 - = # of Instructions of a program * CPI

- CPI?

 - Average Clock Cycles per Instruction






CPU Time (cont.)

= Instr. Count * CPI * Clock Cycle Time

= (Instr. Count * CPI) / Clock Rate

$$\text{Processor Performance} = \frac{\text{Time}}{\text{Program}}$$

<div data-bbox="324 978 774 1160"> The picture can't be displayed. Instructions = $\frac{\text{-----}}{\text{Program}}$</div>	X	<div data-bbox="871 978 1286 1160"> The picture can't be displayed. Cycles $\frac{\text{-----}}{\text{Instruction}}$</div>	X	<div data-bbox="1367 978 1721 1160"> The picture can't be displayed. Time $\frac{\text{-----}}{\text{Cycle}}$</div>
(code size)		(CPI)		(cycle time)



CPU Time (cont.)

- Example
 - Instruction count (IC) = 10,000
 - CPI = 4
 - Clock cycle time = 500ps or 0.5 ns
 - ➔ CPU time = $10,000 * 4 * 0.5 = 20 \text{ us}$
- Question:
 - IC same as code size or lines of code?



CPU Time (cont.)

- How these Parameters Determined?
 - Instruction count
 - CPI
 - Clock cycle time



CPU Time (cont.)


- **Instruction Count (IC)**
 - Determined by program (e.g., Algorithm)
 - ISA
 - Compiler
- **CPI**
 - Determined by uArch
 - The way processor is implemented
 - Code CPI also depends on program
 - Compiler (?)
- **Clock Cycle Time**
 - Determined by uArch and Technology



Easy Practice

- Assume:
 - A C program compiled on two computers
 - Computer A with a RISC ISA
 - Computer B with a CISC ISA
 - Q1: which one would have **higher IC**?
 - Q2: which one would have **smaller CPI**?
 - Q3: which one would have **higher performance**?
 - Q4: can we have **CPI less than one**?

In superscalar architectures, multiple instructions can be issued and executed in a single clock cycle. Out-of-order execution can also lead to lower average CPI, where the CPU rearranges instruction execution for optimal performance. Pipelined processors can potentially execute multiple stages of different instructions simultaneously, which can improve throughput and reduce the overall CPI. However, the CPI for individual instructions is still typically greater than or equal to 1, but the average CPI across a range of instructions in a program can definitely be less than 1, particularly for highly optimized workloads.



CPI Classes

- Question:
 - Do all instructions have same CPI?
 - No

	CPI for instruction classes		
	A	B	C
CPI	1	2	3

Code Sequence	Instruction count		
	A	B	C
1	2	1	2
2	4	1	1



CPI Classes (cont.)

- CPU Clock Cycles

$$= \text{CPI}_1 * \text{IC}_1 + \text{CPI}_2 * \text{IC}_2 + \dots + \text{CPI}_n * \text{IC}_n$$

	CPI for instruction classes		
	A	B	C
CPI	1	2	3

Code Sequence	Instruction count		
	A	B	C
1	2	1	2
2	4	1	1



CPI in More Detail

- If different instruction classes take different numbers of cycles

$$\text{Clock Cycles} = \sum_{i=1}^n (\text{CPI}_i \times \text{Instruction Count}_i)$$

- **Weighted Average CPI**

$$\text{CPI} = \frac{\text{Clock Cycles}}{\text{Instruction Count}} = \sum_{i=1}^n \left(\text{CPI}_i \times \frac{\text{Instruction Count}_i}{\text{Instruction Count}} \right)$$

Relative frequency



Performance Evaluation of Computers

- So Far
 - Learnt how to measure CPU Performance
- Consider Two CPUs
 - CPU1 runs faster than CPU2 on program A
 - CPU2 runs faster than CPU1 on program B
- Questions:
 - How we can decide which CPU is faster?
 - Which candidate programs should we choose to compare CPU1 and CPU2?



Performance Evaluation of Computers (cont.)

- Assume
 - A user typically runs programs A, B, and C in his computer (CPU1)
 - “Workload”
 - Set of programs A, B, and C
 - If CPU2 runs this workload faster
 - ➔ $\text{Performance}(\text{CPU2}) > \text{Performance}(\text{CPU1})$
 - In reality, we use a set of programs called **benchmarks** to compare performance of processors



Performance Evaluation of Computers (cont.)

- Which CPU Runs a Workload Faster?
 - Depends on type of average we take over execution times
 - Simple arithmetic mean
 - Weighted arithmetic mean
 - Average over ratios
 - Geometric mean

CPI	Computer A	Computer B
Program 1	1	10
Program 2	1000	100
Total Time	1001	110



Performance Evaluation of Computers (cont.)

- Another Question:
 - Why not just running one simple program to compare performance of CPUs?
 - Agreeing on one simple program very hard
 - Designers can optimize processors towards fast running that simple program



Performance Evaluation of Computers (cont.)

- Which Programs to Choose?
 - Real programs such as MS-Word, Internet explorer, latex compilers
 - Synthetic benchmarks
 - Emulate frequency of different instructions in real programs
 - Standard benchmarks
 - Examine processor and memory hierarchy



Performance Evaluation of Computers (cont.)

- Which Programs to Choose?
 - Programs from different applications
 - SPEC CPU2000 benchmarks

Integer Benchmarks (12)		FP Benchmarks (14)	
Name	Description	Name	Description
gzip	Compression	ammp	Computation chemistry
vpr	FPGA circuit P&R	swim	Shallow water model
gcc	C compiler	art	Image recognition using neural network
parser	Word processing program	galgel	Computational fluid dynamics




Easy Practice

- Question:
 - Can microprocessor designers optimize their architecture to optimize performance of SPEC2000 or other benchmark program?



Performance of PIII and P4 for SPEC2000



CINT2006 for Intel Core i7 920

Description	Name	Instruction Count x 10 ⁹	CPI	Clock cycle time (seconds x 10 ⁻⁹)	Execution Time (seconds)	Reference Time (seconds)	SPECratio
Interpreted string processing	perl	2252	0.60	0.376	508	9770	19.2
Block-sorting compression	bzip2	2390	0.70	0.376	629	9650	15.4
GNU C compiler	gcc	794	1.20	0.376	358	8050	22.5
Combinatorial optimization	mcf	221	2.66	0.376	221	9120	41.2
Go game (AI)	go	1274	1.10	0.376	527	10490	19.9
Search gene sequence	hmmer	2616	0.60	0.376	590	9330	15.8
Chess game (AI)	sjeng	1948	0.80	0.376	586	12100	20.7
Quantum computer simulation	libquantum	659	0.44	0.376	109	20720	190.0
Video compression	h264avc	3793	0.50	0.376	713	22130	31.0
Discrete event simulation library	omnetpp	367	2.10	0.376	290	6250	21.5
Games/path finding	astar	1250	1.00	0.376	470	7020	14.9
XML parsing	xalancbmk	1045	0.70	0.376	275	6900	25.1
Geometric mean	–	–	–	–	–	–	25.7



Experiment Setup

- How to Report Experiment Setup?
 - CPU frequency is NOT enough!
 - Need to Report both HW & SW config.
- Software
 - Operating system: WinXP prof. SP1
 - Compiler: Microsoft Visual Studio.NET
 - 7.0.xxx
 - File system type: NTFS
 - System state: default
 - Benchmark, Program, Input to program



Experiment Setup (cont.)

- Hardware
 - Hardware vender: Dell
 - Model number: Precision WorkStation 360
 - CPU: Intel Pentium 4 (800 MHz system bus)
 - CPU MHz: 3200
 - FPU: Integrated
 - Primary cache: 12KB (I), 8KB (D), both on-chip
 - Secondary cache: 512KB (I+D), on-chip
 - L3 cache: 2048KB (I+D), on-chip
 - Memory: 4x512MB ECC DDR400 SDRAM CL3
 - Disk subsystem: 1x80GB ATA/100 7200 RPM



Benchmark Classification

- **CPU Intensive** Benchmarks
 - Used to evaluate performance of new micro-architectures
- **Memory Intensive** Benchmarks
 - Used to evaluate configurations of cache/memory hierarchy
- **I/O Intensive** Benchmarks
 - Used to evaluate performance of I/O devices and storage subsystem



CPU/Memory/IO-Intensive Benchmarks

- SPEC Benchmarks
 - CPU/Memory intensive benchmarks
 - Some stress CPU
 - Some stress memory subsystem
- SPEC-Web
 - I/O intensive benchmarks
 - Mostly stress I/O subsystem
 - Disk subsystem, network connections, ...



SPECWeb99 Performance for Variety of Systems

System	Processor	Number of disk drives	Number of CPUs	Number of networks	Clock rate (GHz)	Result
1550/1000	Pentium III	2	2	2	1	2765
1650	Pentium III	3	2	1	1.4	1810
2500	Pentium III	8	2	4	1.13	3435
2550	Pentium III	1	2	1	1.26	1454
2650	Pentium 4 Xeon	5	2	4	3.06	5698
4600	Pentium 4 Xeon	10	2	4	2.2	4615
6400/700	Pentium III Xeon	5	4	4	0.7	4200
6600	Pentium 4 Xeon MP	8	4	8	2	6700
8450/700	Pentium III Xeon	7	8	8	0.7	8001




Speedup

- $\text{Speedup} = \text{time}_{\text{original}} / \text{time}_{\text{improved}}$
- Example
 - $\text{time}_{\text{original}} = 100 \text{ s}$
 - $\text{time}_{\text{improved}} = 98 \text{ s}$
 - $\text{Speedup} = 100/98 = 1.02$



Amdahl's Law

- Amdahl's Law Says
 - Speedup limited to fraction improved
 - Obvious, but fundamental observation



Amdahl's Law (cont.)

- Obvious but Common Mistakes
 - CPU upgraded from 1.5Ghz to 3Ghz
 - But not seeing 100% improvement in performance
 - Combinational upgrade 1
 - CPU from 1.5Ghz to 3Ghz
 - Memory from 1GB-90nm to 2GB-45nm
 - Combinational upgrade 2
 - CPU from 1.5Ghz to 3Ghz
 - Memory from 1GB-90nm to 2GB-45nm
 - L1 from 16KB to 32KB & L2 from 1MB to 2MB



Speedup (cont.)

ExecTime_{new}

$$= \text{ExecTime}_{\text{old}} \times \{(1 - \text{Frac}_{\text{enhanced}}) + (\text{Frac}_{\text{enhanced}} / \text{Speedup}_{\text{enhanced}})\}$$

Speedup_{overall} = $\text{ExecTime}_{\text{old}} / \text{ExecTime}_{\text{new}} =$

1

$$(1 - \text{Frac}_{\text{enhanced}}) + (\text{Frac}_{\text{enhanced}} / \text{Speedup}_{\text{enhanced}})$$



MIPS: Performance Metric

- MIPS
 - Million Instructions Per Second

$$\begin{aligned}\text{MIPS} &= \frac{\text{Instruction count}}{\text{Execution time} \times 10^6} \\ &= \frac{\text{Instruction count}}{\frac{\text{Instruction count} \times \text{CPI}}{\text{Clock rate}}} \times 10^6 = \frac{\text{Clock rate}}{\text{CPI} \times 10^6}\end{aligned}$$

- Drawbacks?



MIPS: Performance Metric (cont.)

- Drawbacks
 - Not taking into account capabilities of instructions
 - Comparing two different ISAs not fair
 - Not realistic metric even on same CPU
 - Two programs: program A and program B
 - $MIPS(A) > MIPS(B) \rightarrow Perf.(A) > Perf.(B) ???$
 - Some optimization tech. add more code



MIPS: Performance Metric (cont.)

- Example:
 - Machine A
 - Special instruction for performing square root
 - It takes 100 cycles to execute
 - Machine B
 - Doesn't have special instruction
 - must perform square root in software using simple instructions
 - e.g, Add, Mult, Shift each take 1 cycle to execute
 - Clock cycle = 1us
 - Machine A: $1/100 \text{ MIPS} = 0.01 \text{ MIPS}$
 - Machine B: 1 MIPS



MFLOPS: Performance Metric

- MFLOPS
 - = (FP ops/program) x (program/time) x 10^{-6}
- Popular in scientific computing
 - FP ops were previously much slower than regular instructions (i.e., off-chip, sequential execution)
- Not great for “predicting” performance
 - Ignores other instructions (e.g., load/store)
 - Not all FP ops have common format
 - Depends on how FP-intensive program is



MFLOPS: Performance Metric (cont.)

- MFLOPS → GFLOPS → TFLOPS → PFLOPS
→ EFLOPS

Rank	System	Cores	Rmax (TFlop/s)	Rpeak (TFlop/s)	Power (kW)
1	Supercomputer Fugaku - Supercomputer Fugaku, A64FX 48C 2.2GHz, Tofu interconnect D, Fujitsu RIKEN Center for Computational Science Japan	7,630,848	442,010.0	537,212.0	29,899
2	Summit - IBM Power System AC922, IBM POWER9 22C 3.07GHz, NVIDIA Volta GV100, Dual-rail Mellanox EDR Infiniband, IBM DOE/SC/Oak Ridge National Laboratory United States	2,414,592	148,600.0	200,794.9	10,096
3	Sierra - IBM Power System AC922, IBM POWER9 22C 3.1GHz, NVIDIA Volta GV100, Dual-rail Mellanox EDR Infiniband, IBM / NVIDIA / Mellanox DOE/NNSA/LLNL United States	1,572,480	94,640.0	125,712.0	7,438
4	Sunway TaihuLight - Sunway MPP, Sunway SW26010 260C 1.45GHz, Sunway, NRCPC National Supercomputing Center in Wuxi China	10,649,600	93,014.6	125,435.9	15,371

- <https://www.top500.org/lists/top500/2021/11/>





Thanks for Your Attention!

